

Document 6 — How to Present and Pitch

SIH26166 · Delivery, division of labor, Q&A preparation, and the claims to make and avoid

Drafted August 31, 2026

Contents

1. The narrative arc that actually lands	2
2. Division of labor for delivery and Q&A	2
3. Tone: declarative, not apologetic	2
4. Claims to make, and claims to avoid	2
5. Anticipated question tiers — prepare, don't improvise	3
6. The single failure mode to avoid at all costs	3
7. The one sentence worth saying out loud, verbatim	3

1. The narrative arc that actually lands

Open with the failure, not your result: *“Here’s where SIFT falls apart on lunar shadows”* — show it visibly failing — then show your fix holding exactly where it broke. This ordering is stronger than presenting your numbers cold, because it gives the judges the problem before the solution, in the order they’ll actually be able to follow and evaluate it.

2. Division of labor for delivery and Q&A

Assign these explicitly **before** you’re in the room — don’t improvise on the day:

- **One person owns the math** — the equation on Slide 5, and any first-principles question about it.
- **One person owns the data/build story** — registration lag, scene selection, what’s real data versus synthetic validation.
- **One person owns scope and honesty** — the “here’s what we didn’t solve, and why” answers.

3. Tone: declarative, not apologetic

Compare these two ways of saying the same fact:

- *“Our shadow check shares a DEM dependency with the geometry it validates; the fix is TMC-2’s own stereo triplet, which we scoped out for time.”* — reads as controlled, deliberate.
- *“We didn’t have time to fully solve that.”* — reads as unfinished, even though it’s the same underlying limitation.

State limitations as design decisions with stated reasons, not as apologies. This single tone shift changes how every honest answer in your Q&A is received.

4. Claims to make, and claims to avoid

Make these claims — they’re accurate and defensible: - “We built the generic pipeline the PS asked for, across OHRC, TMC-2, and IIRS, against external reference imagery — not a single-failure-mode demo.” - “Classical matching is documented (ISRO SAC, 2025) to fail under lunar shadow and subtle intensity shift — that’s the exact failure our shadow-consistency layer targets.” - “Everything in our geometry and photometric pipeline is closed-form and exact; the only trained component is recognizing that two patches under different sun angles show the same ground feature.”

Avoid these claims — they’re unverifiable and one pointed question collapses them: - “This can’t be done by AI” / “no AI could replicate this” — untestable, and invites exactly the question you can’t answer well. - “We’ve beaten NASA/ISRO’s own methods” — you don’t have access to ISRO’s undocumented internal pipeline (OPTIMUS) to make this comparison at all; the accurate claim is narrower: you targeted a *specific, published, named* gap in the public benchmark literature. - Any specific accuracy number you didn’t personally measure on your own scene pair.

5. Anticipated question tiers — prepare, don’t improvise

Tier 1 (expect these, answer instantly): - “What RMSE did you achieve, and on real or synthetic data?” - “Why did you pick this invariance strategy over the alternative?” - “What’s your inlier ratio, and how does it compare to the classical baseline?”

Tier 2 (know cold, may take a beat to answer): - “How do you know your improvement isn’t just overfitting to your one scene pair?” - “What happens with IIRS — is that a real result or a fallback?” - “How do you validate accuracy without ground truth?” (Answer: crater-anchor points + synthetic DEM pairs, cross-checked against each other, per your own Log 04 Open 5 — and be honest that a shared systematic bias could still fool both checks simultaneously.)

Tier 3 (know these exist, be able to name them without a full answer): - Space-weathering drift between imaging epochs. - Ephemeris precision bounds on illumination geometry. - Geodetic/datum reconciliation across ISRO/NASA/JAXA control networks. Naming that a question is a known open item on your own list is a stronger answer than pretending it doesn’t exist, or improvising a confident-sounding but unverified number.

6. The single failure mode to avoid at all costs

Claiming a number you didn’t measure, or implying a real-data result exists when only the synthetic validation actually ran. One caught fabrication undoes every other honest answer you gave in the room — judges remember the one inconsistency far more than the ten correct statements around it. If you’re not sure whether a specific number came from real data or synthetic validation, say so explicitly rather than guessing which one it was live.

7. The one sentence worth saying out loud, verbatim

“We deliberately didn’t ask an AI to guess at results we could only get by running real experiments on real ISRO data — and everywhere we did use AI, we can point to exactly what it accelerated and how we verified it.”

This is true if you’ve actually kept the AI-boundary discipline from your own workplan (Document 1’s Part C validation step, Document 2’s derivation-over-fabrication approach, and never reporting a metric you didn’t compute) — and in a room full of teams who likely leaned on AI less carefully, saying it plainly is a genuine differentiator, not a hedge.